

CS339 lab7 Vxlan

hm1

在VM1和VM2之间搭建好Vxlan后，在VM2中输入

```
h1 ping 10.0.0.1
```

打开Wireshark后获得的监控信息如下：

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
2	0.000557392	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
3	1.023469930	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
4	1.023910027	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
5	2.047882912	e6:6f:a7:00:31:34	06:3d:3f:fc:94:7c	ARP	42	Who has 10.0.0.1?
6	2.048144091	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
7	2.048524263	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
8	2.048880092	06:3d:3f:fc:94:7c	e6:6f:a7:00:31:34	ARP	42	10.0.0.3 is at e6:6f:a7:00:31:34
9	3.071314860	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
10	3.071851612	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
11	4.096002849	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
12	4.096516737	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
13	5.119378373	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
14	5.119787502	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
15	6.143188309	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request
16	6.143655636	10.0.0.3	10.0.0.1	ICMP	98	Echo (ping) reply
17	6.830794965	fe80::a093:6fff:fe9...	ff02::fb	MDNS	180	Standard query
18	7.163750050	10.0.0.1	10.0.0.3	ICMP	98	Echo (ping) request

可以看出，在不同的虚拟机的host之间通信时，有利用到ARP协议来通过MAC地址获得IP地址，再通过ICMP进行传包的，可以看出Vxlan确实是利用MAC-in-UDP把VM1和VM2之间的壁垒给打通了。

hm2

Bandwith between VM1's IPs and VM2's IPs

<pre>sanliu@ubuntu:~\$ iperf -s Server listening on TCP port 5001 TCP window size: 85.3 KByte (default) [4] local 192.168.86.133 port 5001 connected with 192.168.86.128 port 38936 [10] Interval Transfer Bandwidth [4] 0.0-10.0 sec 4.38 Gbytes 3.75 Gbits/sec</pre>		<pre>sanliu@ubuntu:~\$ iperf -c 192.168.86.133 Client connecting to 192.168.86.133, TCP port 5001 TCP window size: 85.3 KByte (default) [3] local 192.168.86.128 port 38936 connected with 192.168.86.133 port 5001 [10] Interval Transfer Bandwidth [3] 0.0-10.0 sec 4.38 Gbytes 3.76 Gbits/sec</pre>
---	---	---

bandwith=3.75Gbits/sec

Bandwith between 10.0.0.1 and 10.0.0.3

由于二者是通过Vxlan通信的，因此相比起正常的ICMP通信，这要多出至少50字节的头部报文，因此在使用iperf测试时需要增大端口的mtu值以避免出现发送字节数大于mtu而丢包的情况。

分别在两个虚拟机上输入：

```
h1 ifconfig h1-eth0 mtu 1600 up
```

再使用iperf测试得到结果如下

```
mininet> h1 iperf -c 10.0.0.1 -M 1400
WARNING: attempt to set TCP maximum segment size to 1400, but got 536
-----
Client connecting to 10.0.0.1, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[  3] local 10.0.0.3 port 48714 connected with 10.0.0.1 port 5001
[ ID] Interval      Transfer    Bandwidth
[  3]  0.0-10.0 sec  2.57 GBytes  2.21 Gbits/sec
```

得到二者间的bandwith为2.21Gbits/sec。

在不断的测试中我发现了一个鉴定iperf成功与否的办法，当iperf没有出现错误时，client和server会同时出现响应，而错误时server端不会出现响应。

原因

我认为原因主要在于使用Vxlan传输数据相比起直接传输数据，多了一步从Vxlan报文头部的信息实现Mac映射到IP的解析过程，并且含有Vxlan的报文的长度更长，携带的有效信息相对更少，因此在传输过程中就有了更多的解码消耗并且传输效率有所下降。

hm3

```
6 --- 192.168.86.133 ping statistics ---
6 20 packets transmitted, 20 received, 0% packet loss, time 19389ms
6 rtt min/avg/max/mdev = 0.355/0.500/1.063/0.181 ms
^ samliu@ubuntu:~$
-- 10.0.0.1 ping statistics ---
23 packets transmitted, 23 received, 0% packet loss, time 22465ms
rtt min/avg/max/mdev = 0.434/0.661/2.137/0.422 ms
mininet> 
```

可以看出在两个虚拟机之间直接ping的latency大概是通过Vxlan连接的两个host之间ping的latency的70%。latency和bandwith的数据成反比，即：

$$latency_{vm} * bandwidth_{vm} \approx latency_{vlan} * bandwidth_{vlan}$$

和预期结果一致。